

Deep Reinforcement Learning

Examples

Examples of Reinforcement Learning


X	O	O
O	X	X
		X

tic-tac-toe board

Tic Tac Toe Game

1. **Goal** - To win the game
2. **State** - Current configuration of the board
3. **Actions** - selecting the cell to place your marker
4. **Reward** - Positive for winning, 0 for tie, negative for loosing

Examples of Reinforcement Learning



Cart-Pole Balancing

- **Goal** — Balance the pole on top of a moving cart
- **State** — Pole angle, angular speed. Cart position, horizontal velocity.
- **Actions** — horizontal force to the cart
- **Reward** — 1 at each time step if the pole is upright

Examples of Reinforcement Learning

Doom*

- **Goal:**
Eliminate all opponents
- **State:**
Raw game pixels of the game
- **Actions:**
Up, Down, Left, Right, Shoot, etc.
- **Reward:**
 - Positive when eliminating an opponent,
negative when the agent is eliminated



** Added for important thought-provoking considerations of AI safety in the context of autonomous weapons systems (see AGI lectures on the topic).*

Comparison with Deep Learning


Types of Learning


- Supervised Learning
- Semi-Supervised Learning
- Unsupervised Learning
- Reinforcement Learning



It's all “supervised” by a loss function!

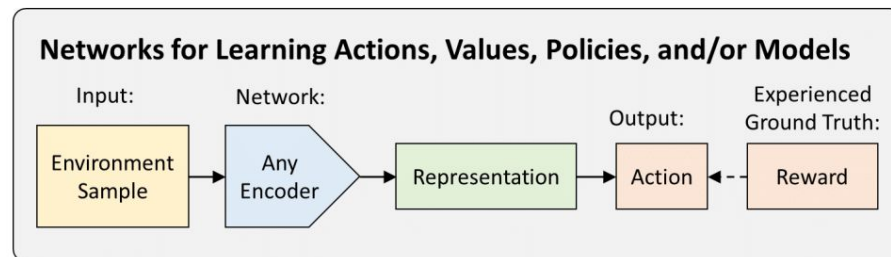
**Someone has to say what's good and what's bad (see Socrates, Epictetus, Kant, Nietzsche, etc.)*

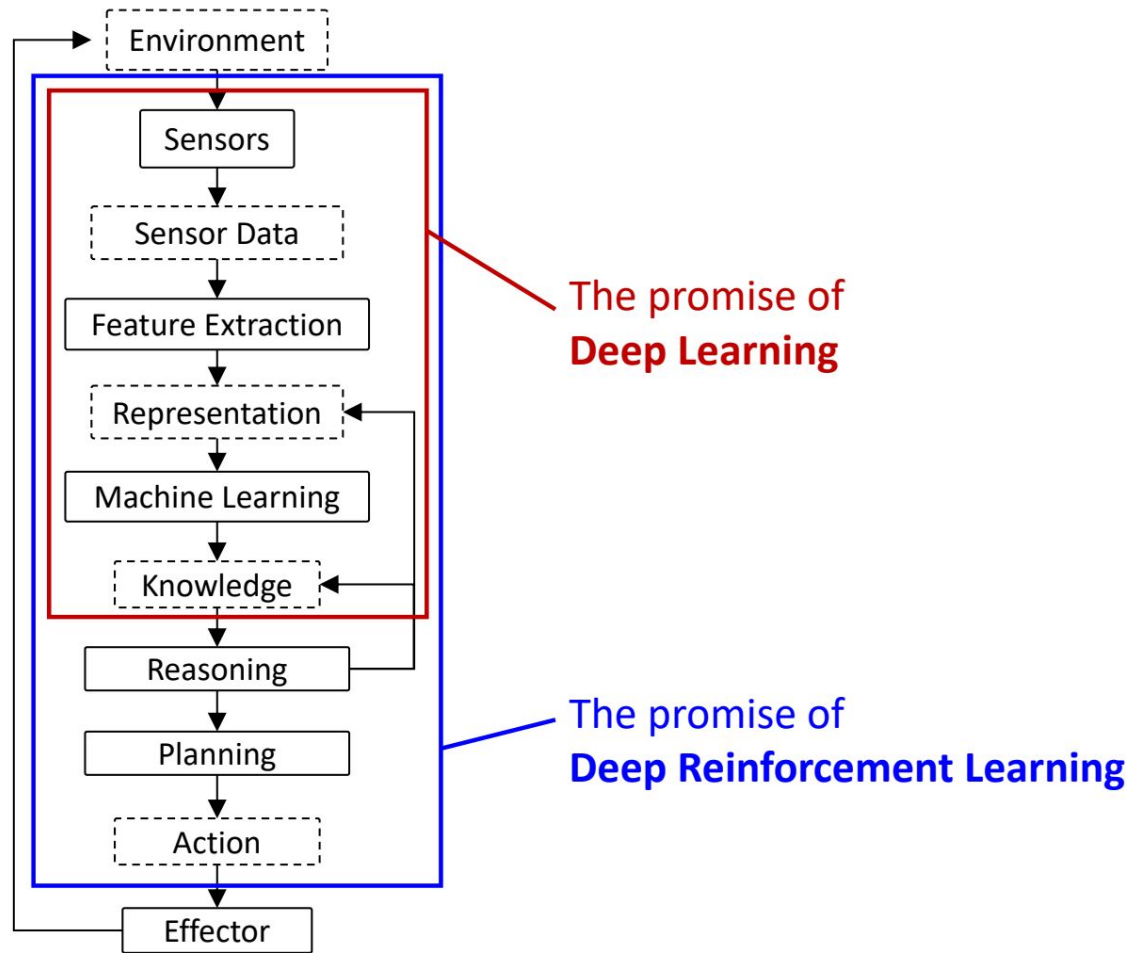




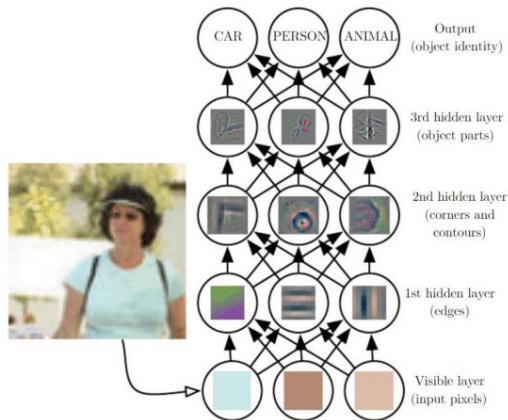

Supervised learning is “teach by **example**”:
Here’s some examples, now learn patterns in these example.

Reinforcement learning is “teach by **experience**”:
Here’s a world, now learn patterns by exploring it.





Deep RL = RL + Neural Networks



Representation Matters

Representation:
The Earth is fixed center of
our Solar System

Representation:
The Sun is fixed center of
our Solar System



Geocentric Model
(Anaximander, 6th century BC)



Heliocentric Model
(Copernicus, 1543)

Background

Value Functions

- The **value of a state** is the expected return starting from that state; depends on the agent's policy:

State - value function for policy π :

$$V^{\pi}(s) = E_{\pi} \left\{ R_t \mid s_t = s \right\} = E_{\pi} \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s \right\}$$


- The **value of taking an action in a state under policy π** is the expected return starting from that state, taking that action, and thereafter following π :

Action - value function for policy π :

$$Q^{\pi}(s, a) = E_{\pi} \left\{ R_t \mid s_t = s, a_t = a \right\} = E_{\pi} \left\{ \sum_{k=0}^{\infty} \gamma^k r_{t+k+1} \mid s_t = s, a_t = a \right\}$$

Q-Learning

- State-action value function: $Q^\pi(s,a)$
 - Expected return when starting in s , performing a , and following π



- Q-Learning: Use **any policy** to estimate Q that maximizes future reward:
 - Q directly approximates Q^* (Bellman optimality equation)
 - Independent of the policy being followed
 - Only requirement: keep updating each (s,a) pair

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha \left(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t) \right)$$

Diagram illustrating the Q-Learning update equation:

- Learning Rate** (α) points to the α term.
- Discount Factor** (γ) points to the γ term.
- New State** (s_t) points to the s_t in $Q_t(s_t, a_t)$ and $Q_{t+1}(s_t, a_t)$.
- Old State** (s_t) points to the s_t in $Q_t(s_t, a_t)$.
- Reward** (R_{t+1}) points to the R_{t+1} term.

Q-Learning: Value Iteration

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha \left(R_{t+1} + \gamma \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t) \right)$$

Diagram illustrating the Q-Learning update equation:


- New State** points to s_t in $Q_{t+1}(s_t, a_t)$.
- Old State** points to s_t in $Q_t(s_t, a_t)$.
- Reward** points to R_{t+1} .
- Learning Rate** points to α .
- Discount Factor** points to γ .

	A1	A2	A3	A4
S1	+1	+2	-1	0
S2	+2	0	+1	-2
S3	-1	+1	0	-2
S4	-2	0	+1	+1

```
initialize  $Q[num\_states, num\_actions]$  arbitrarily
observe initial state  $s$ 
repeat
    select and carry out an action  $a$ 
    observe reward  $r$  and new state  $s'$ 
     $Q[s, a] = Q[s, a] + \alpha(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
     $s = s'$ 
until terminated
```

Q-Learning: Representation Matters

- In practice, Value Iteration is impractical
 - Very limited states/actions
 - Cannot generalize to unobserved states



- Think about the **Breakout** game

- State: screen pixels
 - Image size: **84 × 84** (resized)
 - Consecutive **4** images
 - Grayscale with **256** gray levels




$256^{84 \times 84 \times 4}$ rows in the Q-table!

$= 10^{69,970} \gg 10^{82}$ atoms in the universe

DQN: Deep Q-Learning

Use a neural network to approximate the Q-function:

$$Q(s, a; \theta) \approx Q^*(s, a)$$



Loss Function in Deep Q-Learning

$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[(y_i - Q(s, a; \theta_i))^2 \right],$$

$$y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$$

$$\nabla_{\theta_i} L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot); s' \sim \mathcal{E}} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) - Q(s, a; \theta_i) \right) \nabla_{\theta_i} Q(s, a; \theta_i) \right].$$

DQN and Double DQN

- Loss function (squared error):

$$L = \mathbb{E}[\underbrace{(\mathbf{r} + \gamma \max_{a'} Q(s', a'))}_{\text{target}} - \underbrace{Q(s, a)}_{\text{prediction}}]^2]$$

- DQN: same network for both Q
- Double DQN: separate network for each Q
 - Helps reduce bias introduced by the inaccuracies of Q network at the beginning of training

DQN Tricks

- Experience Replay
 - Stores experiences (actions, state transitions, and rewards) and creates mini-batches from them for the training process
- Fixed Target Network
 - Error calculation includes the target function depends on network parameters and thus changes quickly. Updating it only every 1,000 steps increases stability of training process.

$$Q(s_t, a) \leftarrow Q(s_t, a) + \alpha \left[r_{t+1} + \gamma \max_p Q(s_{t+1}, p) - Q(s_t, a) \right]$$

target Q function in the red rectangular is fixed

Replay	○	○	×	×
Target	○	×	○	×
Breakout	316.8	240.7	10.2	3.2
River Raid	7446.6	4102.8	2867.7	1453.0
Seaquest	2894.4	822.6	1003.0	275.8
Space Invaders	1088.9	826.3	373.2	302.0

Algorithm 1 Deep Q-learning with Experience Replay

Initialize replay memory \mathcal{D} to capacity N

Initialize action-value function Q with random weights

for episode = 1, M **do**

 Initialise sequence $s_1 = \{x_1\}$ and preprocessed sequenced $\phi_1 = \phi(s_1)$

for $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in \mathcal{D}

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from \mathcal{D}

 Set $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$


 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ according to equation 3

end for

end for

Policy Gradient Methods

3 Types of Reinforcement Learning



Model-based

- Learn the model of the world, then plan using the model
- Update model often
- Re-plan often

Value-based


- Learn the state or state-action value
- Act by choosing best action in state
- Exploration is a necessary add-on

Policy-based

- Learn the stochastic policy function that maps state to action
- Act by sampling policy
- Exploration is baked in

Policy Gradient (PG)

- **DQN (off-policy):** Approximate Q and infer optimal policy
- **PG (on-policy):** Directly optimize policy space



Policy Network

Good illustrative explanation:

<http://karpathy.github.io/2016/05/31/rl/>

*"Deep Reinforcement Learning:
Pong from Pixels"*

Policy Objective Functions

- ▶ Goal: given **policy $\pi_{\theta}(s, a)$** , find best **parameters θ**
- ▶ How do we measure the quality of a policy π_{θ} ?
- ▶ In episodic environments we can use the **average total return per episode**
- ▶ In continuing environments we can use the **average reward per step**

Policy Objective Functions: Episodic

► **Episodic-return objective:**

$$\begin{aligned} J_G(\boldsymbol{\theta}) &= \mathbb{E}_{S_0 \sim d_0, \pi_{\boldsymbol{\theta}}} \left[\sum_{t=0}^{\infty} \gamma^t R_{t+1} \right] \\ &= \mathbb{E}_{S_0 \sim d_0, \pi_{\boldsymbol{\theta}}} [G_0] \\ &= \mathbb{E}_{S_0 \sim d_0} [\mathbb{E}_{\pi_{\boldsymbol{\theta}}} [G_t \mid S_t = S_0]] \\ &= \mathbb{E}_{S_0 \sim d_0} [v_{\pi_{\boldsymbol{\theta}}}(S_0)] \end{aligned}$$

where d_0 is the start-state distribution This objective equals the expected value of the start state

Policy Objective Functions: Average Reward

► Average-reward objective

$$\begin{aligned} J_R(\theta) &= \mathbb{E}_{\pi_\theta} [R_{t+1}] \\ &= \mathbb{E}_{S_t \sim d_{\pi_\theta}} \left[\mathbb{E}_{A_t \sim \pi_\theta(S_t)} [R_{t+1} \mid S_t] \right] \\ &= \sum_s d_{\pi_\theta}(s) \sum_a \pi_\theta(s, a) \sum_r p(r \mid s, a) r \end{aligned}$$

where $d_\pi(s) = p(S_t = s \mid \pi)$ is the probability of being in state s in the long run
Think of it as the ratio of time spent in s under policy π

Policy Optimisation

- ▶ Policy based reinforcement learning is an **optimization** problem
- ▶ Find θ that maximises $J(\theta)$
- ▶ We will focus on **stochastic gradient ascent**, which is often quite efficient (and easy to use with deep nets)

Policy Gradient



- ▶ Idea: ascent the gradient of the objective $J(\theta)$

$$\Delta\theta = \alpha \nabla_{\theta} J(\theta)$$

- ▶ Where $\nabla_{\theta} J(\theta)$ is the **policy gradient**

$$\nabla_{\theta} J(\theta) = \begin{pmatrix} \frac{\partial J(\theta)}{\partial \theta_1} \\ \vdots \\ \frac{\partial J(\theta)}{\partial \theta_n} \end{pmatrix}$$

- ▶ and α is a step-size parameter
- ▶ Stochastic policies help ensure $J(\theta)$ is smooth (typically/mostly)



Gradients on parameterized policies

- ▶ How to compute this gradient $\nabla_{\theta} J(\theta)$?
- ▶ Assume policy π_{θ} is differentiable almost everywhere (e.g., neural net)
- ▶ For average reward

$$\nabla_{\theta} J(\theta) = \nabla_{\theta} \mathbb{E}_{\pi_{\theta}}[R] .$$

- ▶ How does $\mathbb{E}[R]$ depend on θ ?

Contextual Bandits Policy Gradient

- ▶ Consider a one-step case (a contextual bandit) such that $J(\theta) = \mathbb{E}_{\pi_\theta}[R(S, A)]$.
(Expectation is over d (states) and π (actions))
(For now, d does **not** depend on π)
- ▶ We cannot sample R_{t+1} and then take a gradient:
 R_{t+1} is just a number and does not depend on θ !
- ▶ Instead, we use the identity:

$$\nabla_\theta \mathbb{E}_{\pi_\theta}[R(S, A)] = \mathbb{E}_{\pi_\theta}[R(S, A) \nabla_\theta \log \pi(A|S)] .$$

(Proof on next slide)

- ▶ The right-hand side gives an expected gradient that can be sampled
- ▶ Also known as REINFORCE (Williams, 1992)

The score function trick

Let $r_{sa} = \mathbb{E}[R(S, A) \mid S = s, A = s]$

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} \mathbb{E}_{\pi_{\boldsymbol{\theta}}}[R(S, A)] &= \nabla_{\boldsymbol{\theta}} \sum_s d(s) \sum_a \pi_{\boldsymbol{\theta}}(a|s) r_{sa} \\&= \sum_s d(s) \sum_a r_{sa} \nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(a|s) \\&= \sum_s d(s) \sum_a r_{sa} \pi_{\boldsymbol{\theta}}(a|s) \frac{\nabla_{\boldsymbol{\theta}} \pi_{\boldsymbol{\theta}}(a|s)}{\pi_{\boldsymbol{\theta}}(a|s)} \\&= \sum_s d(s) \sum_a \pi_{\boldsymbol{\theta}}(a|s) r_{sa} \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(a|s) \\&= \mathbb{E}_{d, \pi_{\boldsymbol{\theta}}}[R(S, A) \nabla_{\boldsymbol{\theta}} \log \pi_{\boldsymbol{\theta}}(A|S)]\end{aligned}$$

Contextual Bandit Policy Gradient

$$\nabla_{\theta} \mathbb{E}[R(S, A)] = \mathbb{E}[\nabla_{\theta} \log \pi_{\theta}(A|S) R(S, A)] \quad (\text{see previous slide})$$

- ▶ This is something we **can** sample
- ▶ Our stochastic policy-gradient update is then

$$\theta_{t+1} = \theta_t + \alpha R_{t+1} \nabla_{\theta} \log \pi_{\theta_t}(A_t | S_t).$$

- ▶ In expectation, this is the following the actual gradient
- ▶ So this is a pure (unbiased) stochastic gradient algorithm
- ▶ Intuition: increase probability for actions with high rewards

Policy Gradient Theorem

- ▶ The policy gradient approach also applies to (multi-step) MDPs
- ▶ Replaces reward R with long-term return G_t or value $q_\pi(s, a)$
- ▶ There are actually two policy gradient theorems (Sutton et al., 2000):
 average return per episode & **average reward per step**

Policy gradient theorem (episodic)

Theorem

For any differentiable policy $\pi_{\theta}(s, a)$, let d_0 be the starting distribution over states in which we begin an episode. Then, the policy gradient of $J(\theta) = \mathbb{E}[G_0 \mid S_0 \sim d_0]$ is

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} \left[\sum_{t=0}^T \gamma^t q_{\pi_{\theta}}(S_t, A_t) \nabla_{\theta} \log \pi_{\theta}(A_t | S_t) \mid S_0 \sim d_0 \right]$$

where

$$\begin{aligned} q_{\pi}(s, a) &= \mathbb{E}_{\pi}[G_t \mid S_t = s, A_t = a] \\ &= \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) \mid S_t = s, A_t = a] \end{aligned}$$

Episodic policy gradients: proof

- Consider trajectory $\tau = S_0, A_0, R_1, S_1, A_1, R_1, S_2, \dots$ with return $G(\tau)$

$$\nabla_{\theta} J_{\theta}(\pi) = \nabla_{\theta} \mathbb{E}[G(\tau)] = \mathbb{E}[G(\tau) \nabla_{\theta} \log p(\tau)] \quad (\text{score function trick})$$

$$\begin{aligned} \nabla_{\theta} \log p(\tau) &= \nabla_{\theta} \log \left[p(S_0) \pi(A_0|S_0) p(S_1|S_0, A_0) \pi(A_1|S_1) \cdots \right] \\ &= \nabla_{\theta} \left[\log p(S_0) + \log \pi(A_0|S_0) + \log p(S_1|S_0, A_0) + \log \pi(A_1|S_1) + \cdots \right] \\ &= \nabla_{\theta} \left[\log \pi(A_0|S_0) + \log \pi(A_1|S_1) + \cdots \right] \end{aligned}$$

So:

$$\nabla_{\theta} J_{\theta}(\pi) = \mathbb{E}_{\pi} [G(\tau) \nabla_{\theta} \sum_{t=0}^T \log \pi(A_t|S_t)]$$

Episodic policy gradients: proof (continued)

$$\begin{aligned}\nabla_{\boldsymbol{\theta}} J_{\boldsymbol{\theta}}(\pi) &= \mathbb{E}_{\pi} \left[G(\tau) \sum_{t=0}^T \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T G(\tau) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\sum_{k=0}^T \gamma^k R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\sum_{k=t}^T \gamma^k R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T \left(\gamma^t \sum_{k=t}^T \gamma^{k-t} R_{k+1} \right) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \\&= \mathbb{E}_{\pi} \left[\sum_{t=0}^T (\gamma^t G_t) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right] \qquad = \mathbb{E}_{\pi} \left[\sum_{t=0}^T \gamma^t q_{\pi}(S_t, A_t) \nabla_{\boldsymbol{\theta}} \log \pi(A_t | S_t) \right]\end{aligned}$$

Episodic policy gradients algorithm

$$\nabla_{\theta} J_{\theta}(\pi) = \mathbb{E}_{\pi} \left[\sum_{t=0}^T \gamma^t q_{\pi}(S_t, A_t) \nabla_{\theta} \log \pi(A_t | S_t) \right]$$

- ▶ We can sample this, given a whole episode
- ▶ Typically, people pull out the sum, and split up this into separate gradients, e.g.,

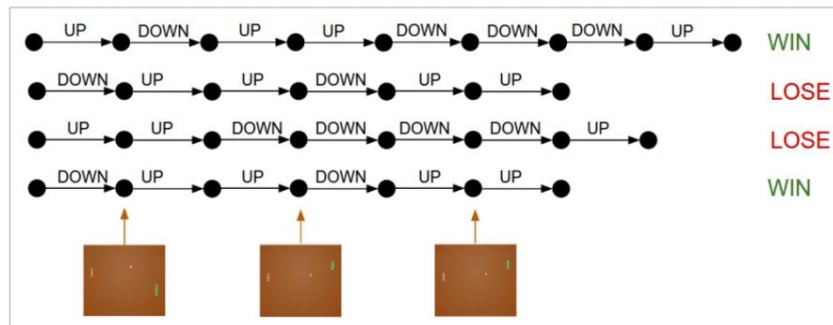
$$\Delta \theta_t = \gamma^t G_t \nabla_{\theta} \log \pi(A_t | S_t)$$

such that $\mathbb{E}_{\pi} [\sum_t \Delta \theta_t] = \nabla_{\theta} J_{\theta}(\pi)$

- ▶ Typically, people ignore the γ^t term, use $\Delta \theta_t = G_t \nabla_{\theta} \log \pi(A_t | S_t)$
- ▶ This is actually okay-ish — we just partially pretend on each step that we could have started an episode in that state instead (alternatively, view it as a slightly biased gradient)

Policy Gradient – Training

Policy Gradients: Run a policy for a while. See what actions led to high rewards. Increase their probability.



- **REINFORCE:** Policy gradient that increases probability of good actions and decreases probability of bad action:

$$\nabla_{\theta} E[R_t] = E[\nabla_{\theta} \log P(a) R_t]$$